


# Efficient Processing of Bursty Information Streams with



TIM O'SHEA ([TIM.OSHEA753@GMAIL.COM](mailto:TIM.OSHEA753@GMAIL.COM) // TIM` @ IRC.FREENODE.NET)

RESEARCH ASSOCIATE, VIRGINIA TECH, ARLINGTON, VA  
PRINCIPAL, O'SHEA RESEARCH LLC



# Rapid Overview

- ▶ An Overview of GNU Radio's Processing Model
  - ▶ Brief primer on the Stream Tag System
  - ▶ Brief primer on GNU Radio's Polymorphic Type
  - ▶ The Concept of a PMT PDU
- ▶ Burst System Motivation
- ▶ Tagged Stream Block (TSB) Based Burst Design
- ▶ Message Port (PDU) Based Burst Design
- ▶ Burst Design Performance Measurement (gr-chunky)
- ▶ Translation Between Modem Design Segments
- ▶ FAß Working Group
- ▶ PyBOMBS Usage Statistics
- ▶ Performance stats: [stats.gnuradio.org](http://stats.gnuradio.org)
- ▶ Rapid GPU Blocks with gr-theano

# GNU Radio Stream Processing

- ▶ We all know and love this
- ▶ It runs forever, the only notion of an “Item” is a “float32” (n bytes)



Standard GNU Radio Stream  
Port with sizeof(float) Item size  
GR Circular Buffer Impl

# GNU Radio Stream Tags

- ▶ Stream tags allow us to annotate events occurring in a stream at precise sample times
- ▶ Great for time tagging, synchronously responding to events in streams, etc

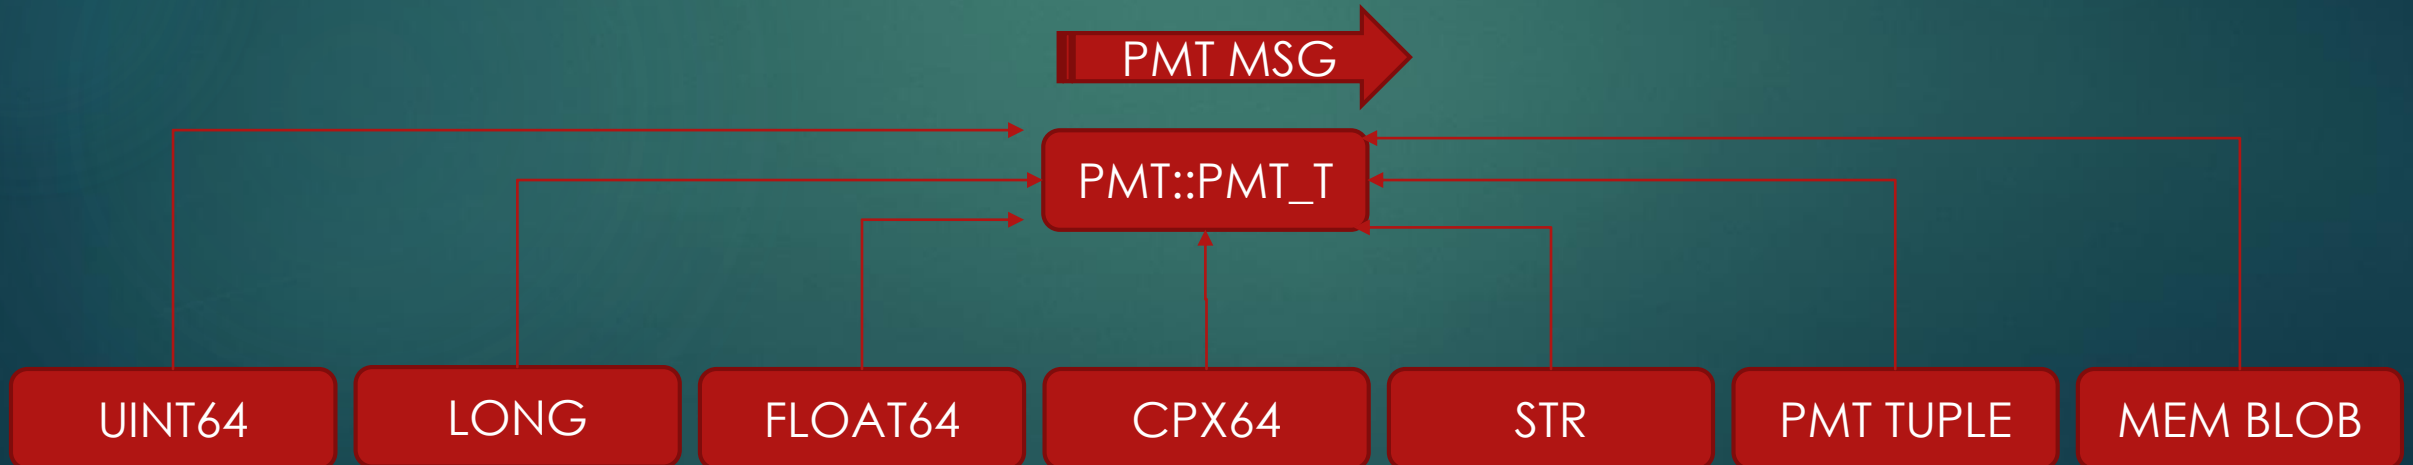


rx\_freq = 2517.750 MHz  
@ sample 8

rx\_time = 9/18/2014 @ 08:38:15.372810  
rx\_freq = 2409.1743 MHz  
@ sample 0

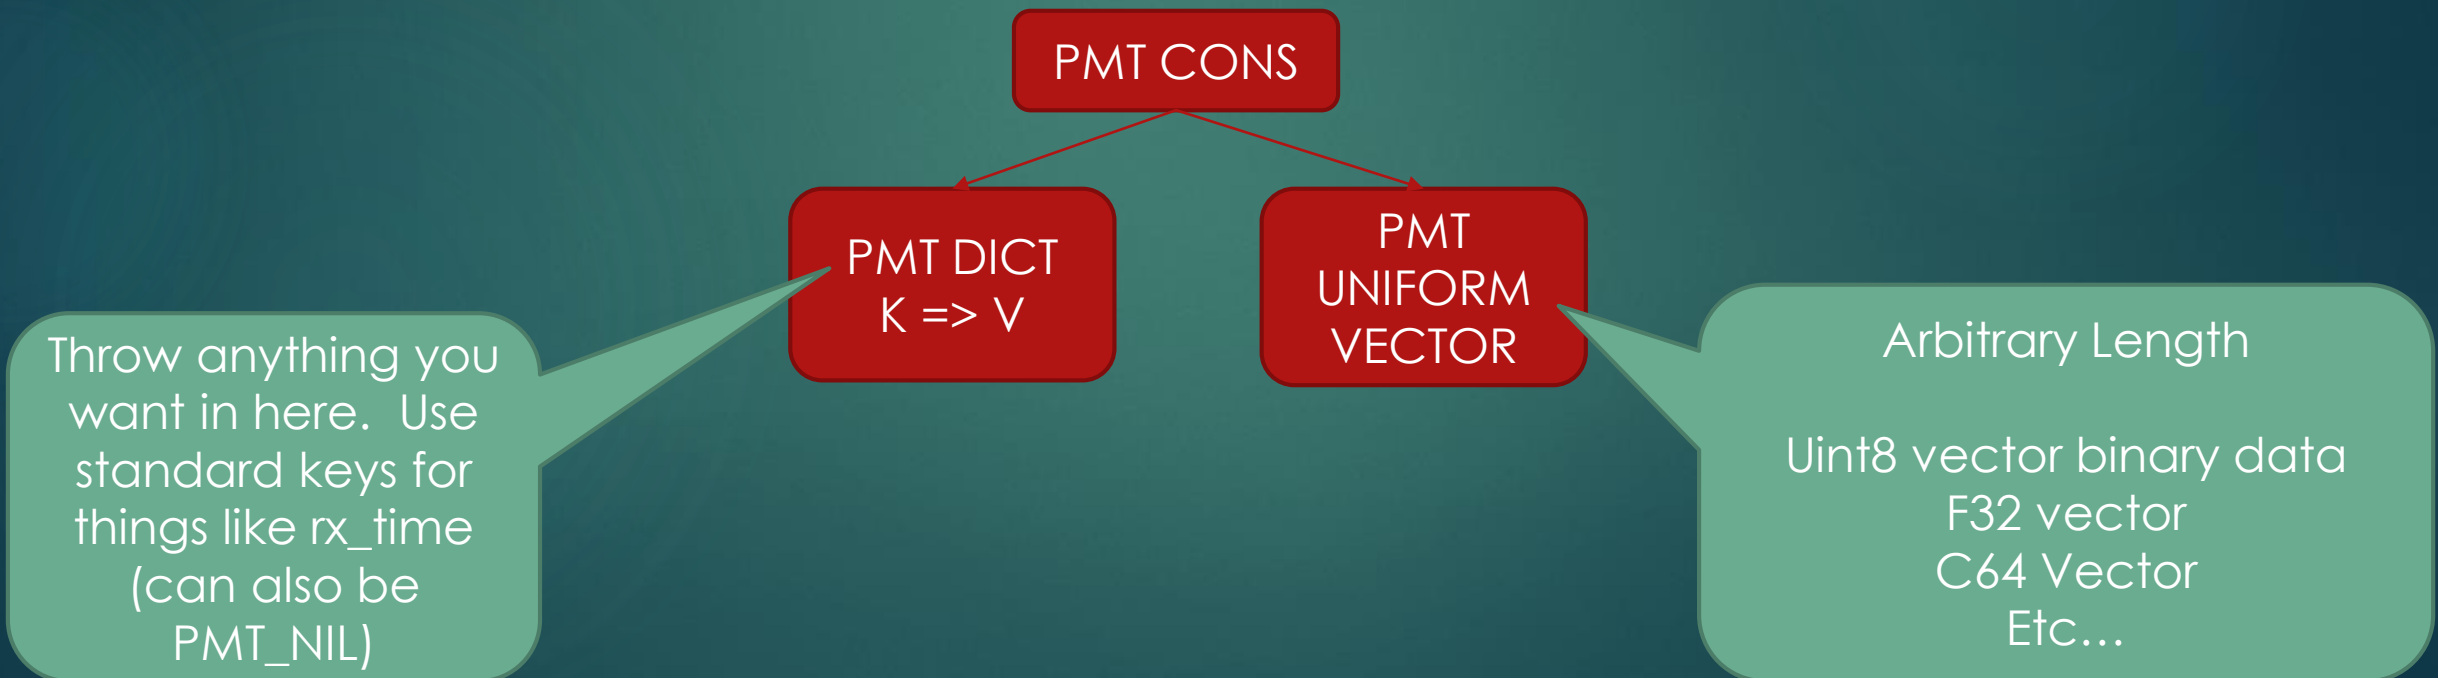
# GNU Radio Polymorphic Type (PMT)

- ▶ Functional Programming in C++ ? Why Not
- ▶ Facilitates Safe Message Passing through several concepts
  - ▶ Everything is a PMT :: OO Inheritance
  - ▶ PMT's are typically write-once constructs – avoid threading hazards
  - ▶ PMT Lifecycles are managed by Boost Shared Pointer ref counting
- ▶ Some PMT Examples



# The PMT Protocol Data Unit (PDU)

- ▶ A “Standard” Interface to carry information between blocks
- ▶ A PDU is a Cons with standard Car and Cdr formats
- ▶ Car = (pmt::dict or pmt::PMT\_NIL), Cdr = (pmt::uvector or list(pmt::uvector))
- ▶ A brief story about underdefined message types (i.e. SCA configure/query strings)
- ▶ Need to start a “standard” keys wiki page on GR site? Uhd has defined some “Standards” already



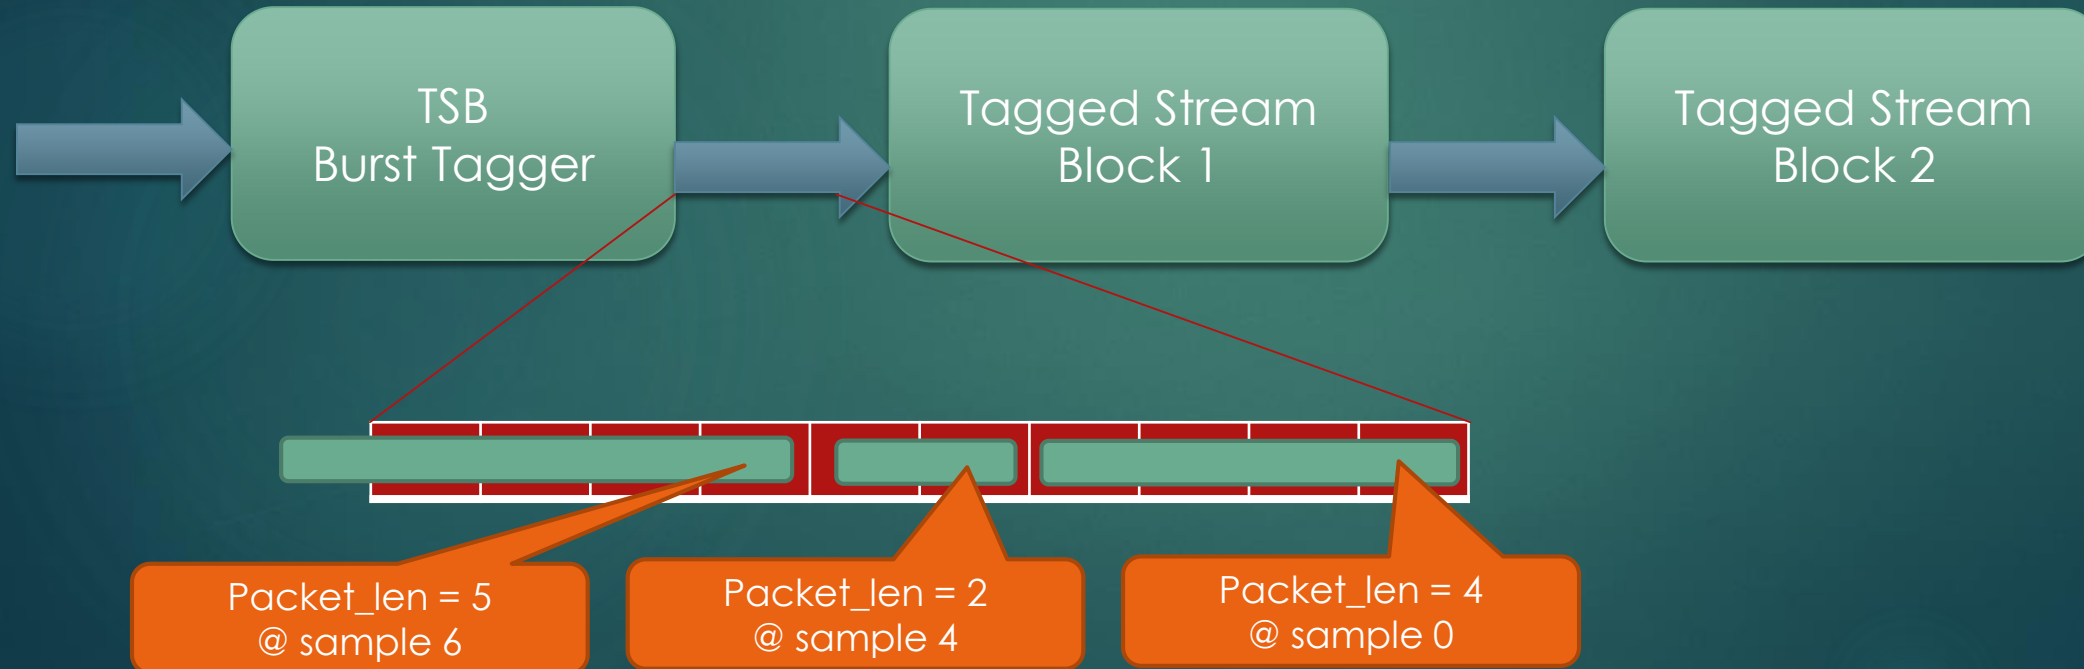
# Motivation for Burst System Design

- ▶ Most systems today are burst or packet based
  - ▶ Systems perform packet based multi-user slot / burst assignment
  - ▶ Synchronous reconfiguration required in many systems
- ▶ Stopping and starting flow graphs is generally not a good way to reconfigure synchronously
- ▶ Using switch blocks typically becomes a nightmare quickly
- ▶ Monolithic stream blocks with lots of internal state are not a great solution
  - ▶ Can be efficient, but generally sacrifice code-reuse & portability



# GNU Radio Tagged Stream Block (TSB)

- ▶ Traditional streaming mode buffers can be used in conjunction with stream tags to pass “bursts” around between blocks intended to process these “bursts”
- ▶ This has the nice side effect that you can often pass them through existing stream blocks without any modification
- ▶ The scheduler is largely unaware this is going on which makes it sad





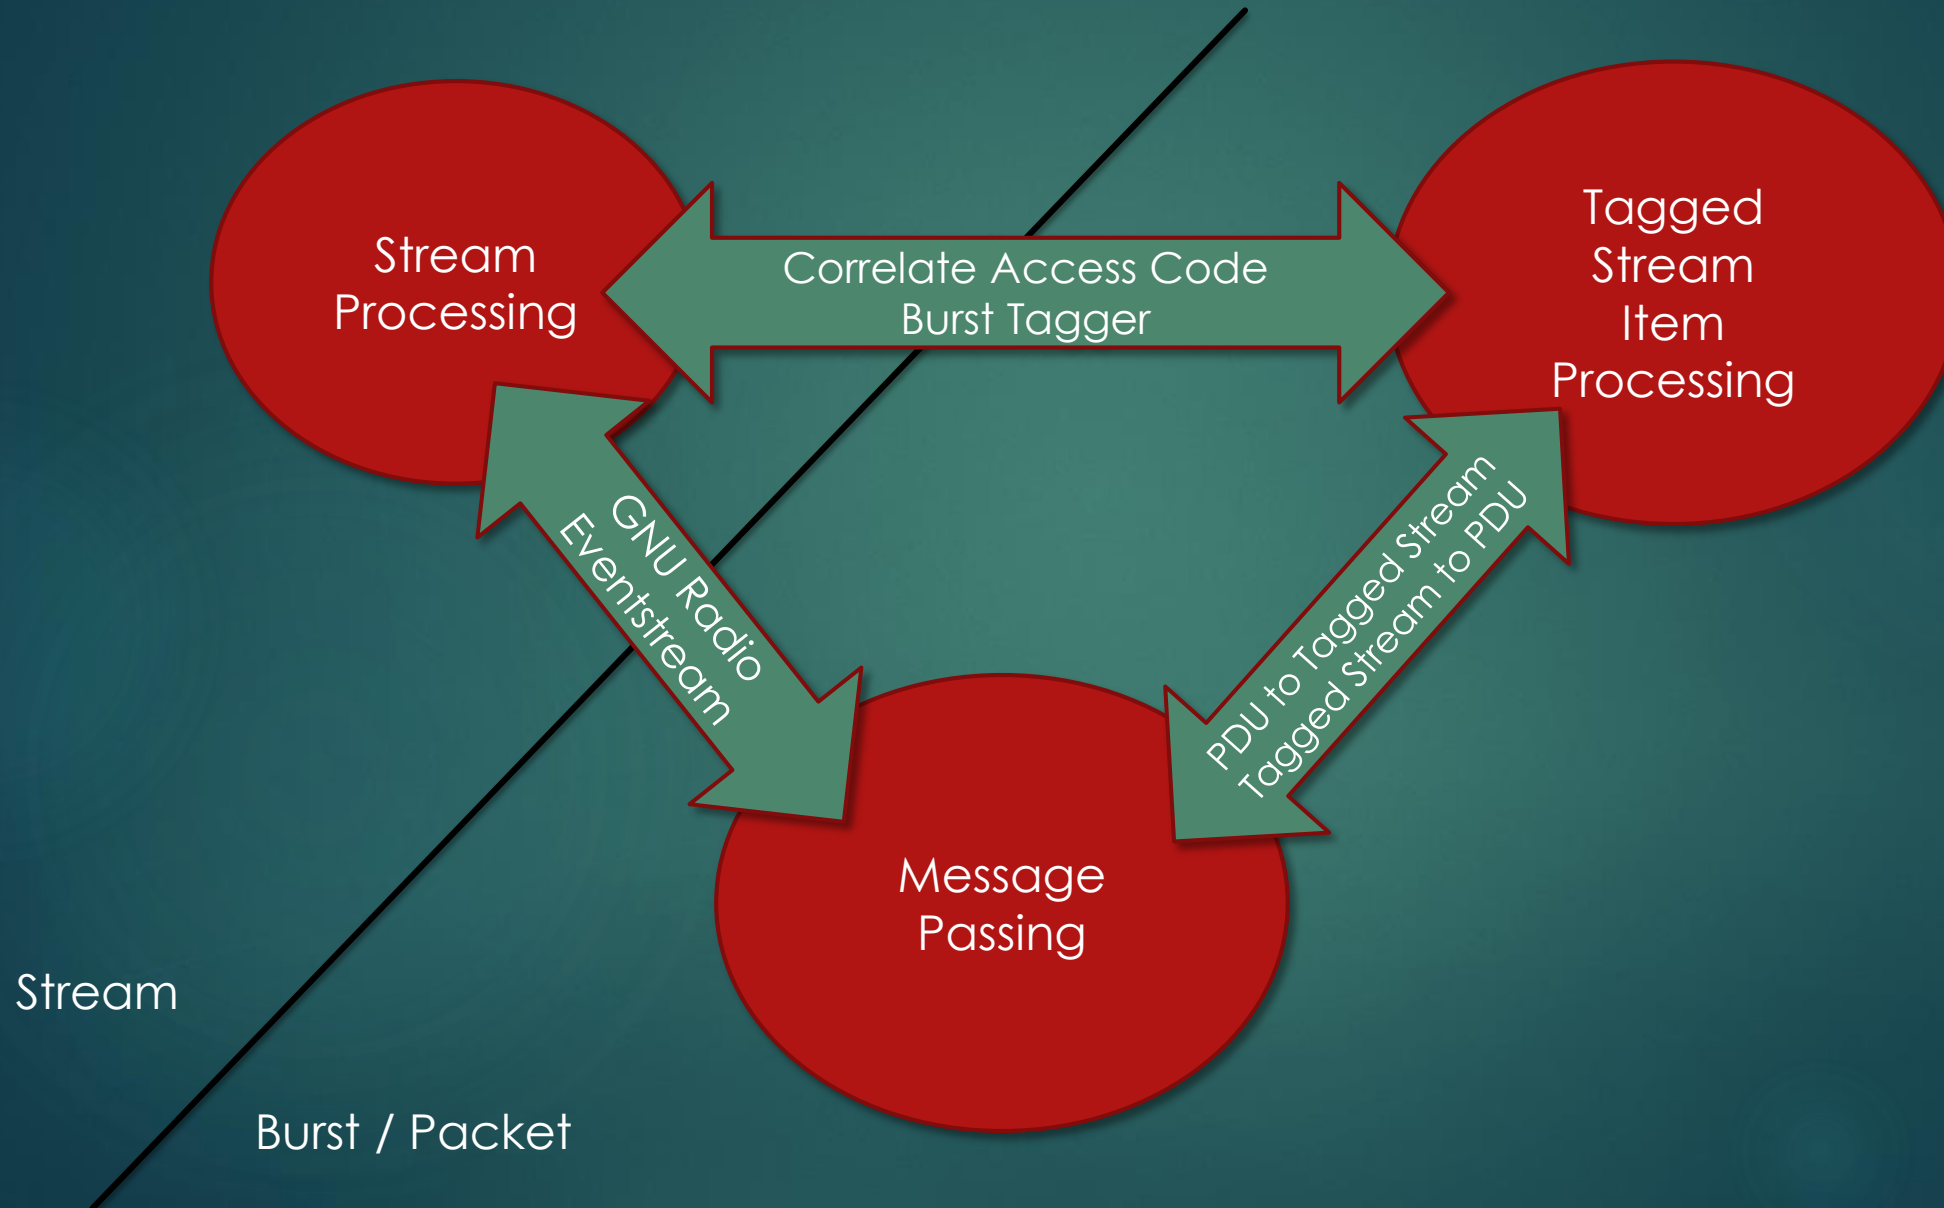
# GNU Radio Message Connections (PDU)

- ▶ Messages don't have to be PDU format, but that is a handy standard
- ▶ Abandon stream ports entirely!!
- ▶ Move to message passing model for “burst” portions of a modem



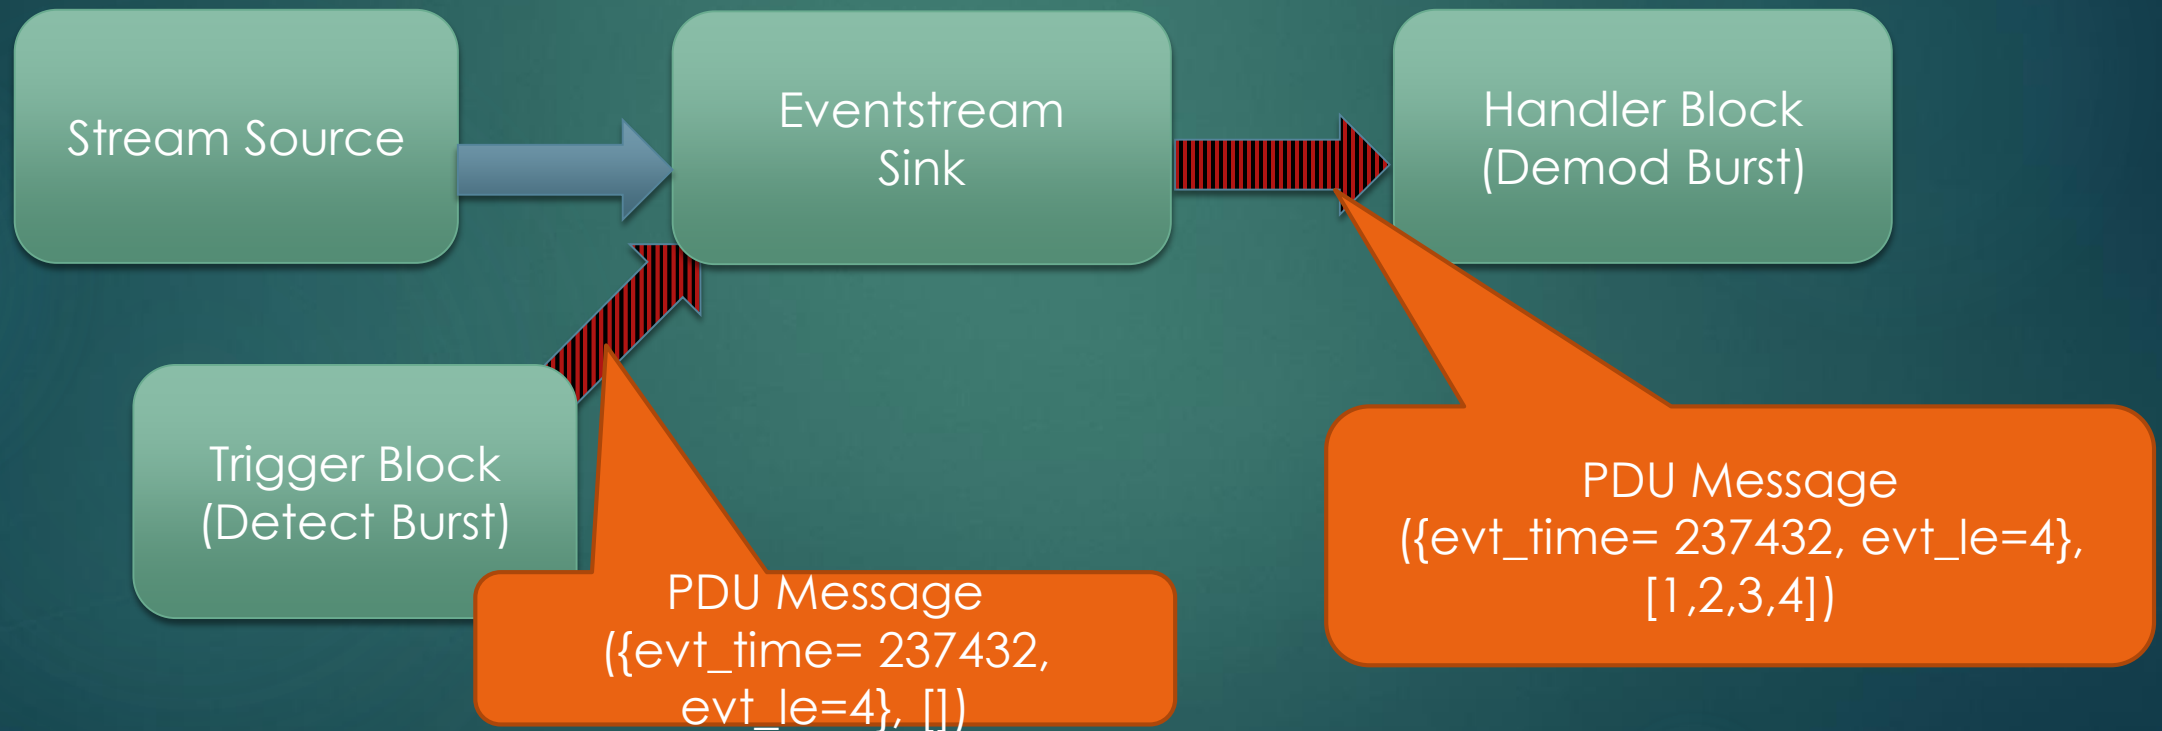
PDU Messages passed between  
block message ports  
Can be arbitrary size, scheduler  
always provides 1 “message”

# Domain Model Translation Blocks



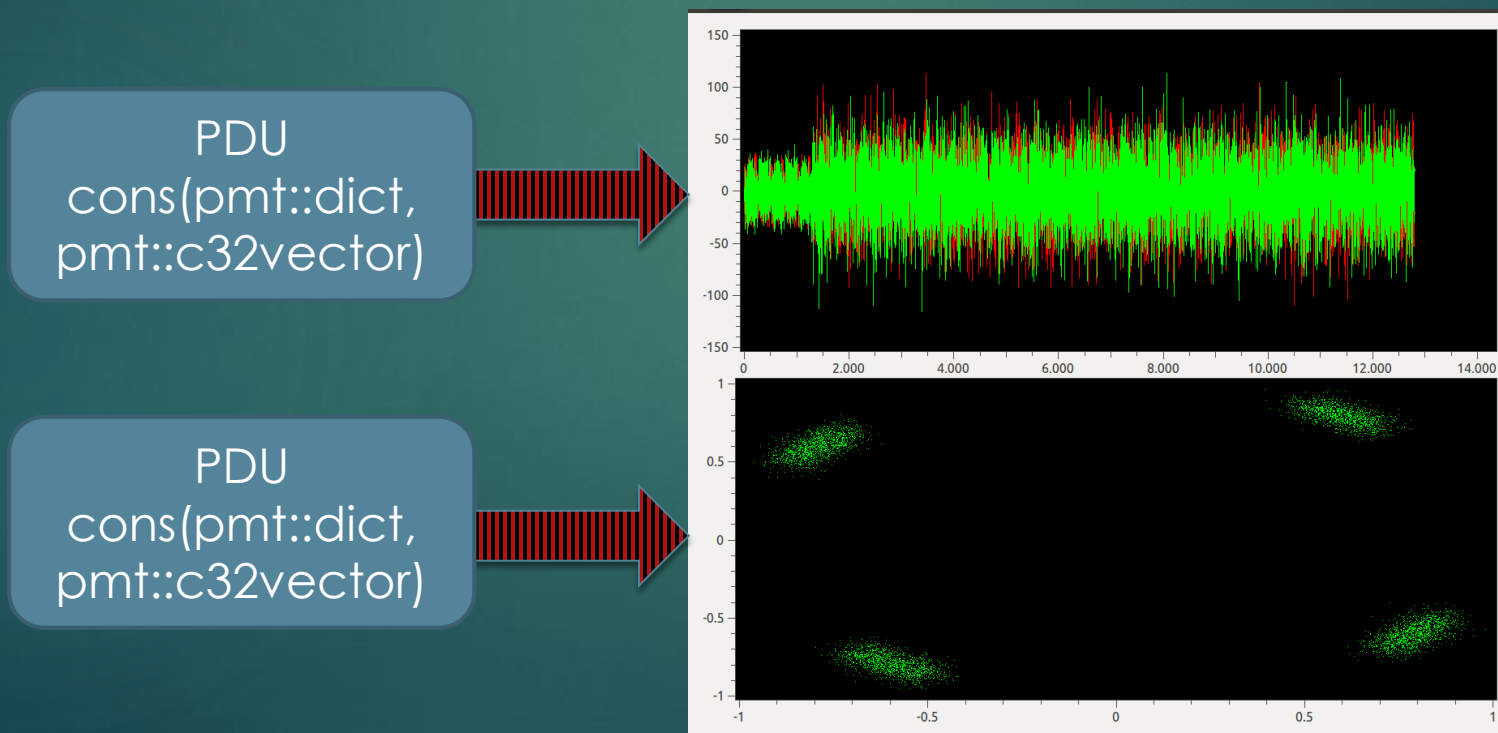
# gr-eventsteam: a brief mention again

- ▶ Provides stream to and from message direct translation
- ▶ Also on github @ <https://github.com/osh/gr-eventstream>
- ▶ Also provides efficient model for concurrent handler execution



# PyQT Burst Plotting Tools (OOT Module)

- ▶ To enable debugging of message based burst applications
- ▶ Plots one message vector at a time instead of N samples at a time
- ▶ Allows you to view one burst, one ofdm symbol, etc at a time
- ▶ Available in pybombs and on github @ <https://github.com/osh/gr-pyqt>



# Functional Architectures for Signal Processing Systems (FAß)

- ▶ Focus on Modem Design Methods in GNU Radio
  - ▶ Specifically TSB & PDU based burst design
- ▶ Ensure we have the tools necessary to build first class systems
- ▶ Improve interoperability and blocks for burst modems
- ▶ Find and fix existing problems in TSB and PDU systems
- ▶ Work on PDU and TSB performance improvement and characterization
- ▶ Enable awesome dynamic systems!
- ▶ Stick around for the WG shortly after this talk to discuss this!

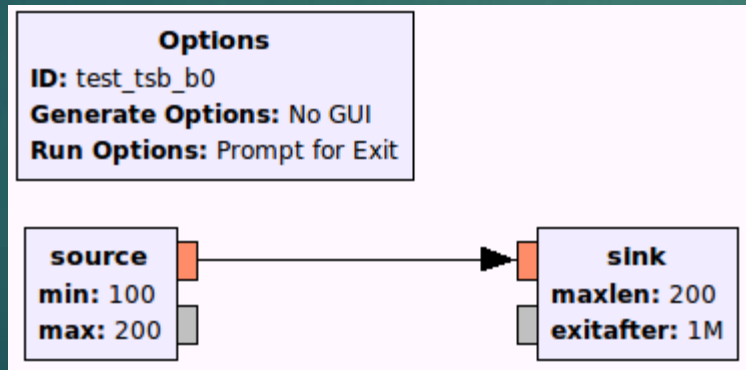


# TSB / PDU Design Pros and Cons

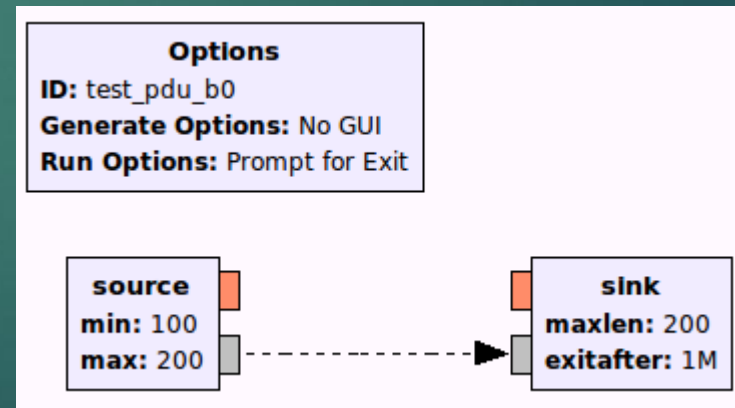
	Pros	Cons	Mitigations
Tagged Stream Blocks	<ul style="list-style-type: none"><li>• Interoperability / Existing Blocks</li><li>• Re-use efficient circular buffers</li></ul>	<ul style="list-style-type: none"><li>• Scheduler not “burst” aware</li><li>• Buffer imposed size limitations</li></ul>	<ul style="list-style-type: none"><li>• Scheduler Improvements?</li></ul>
Message Passing	<ul style="list-style-type: none"><li>• Simpler Implementation</li><li>• Scheduler is burst-size agonistic</li></ul>	<ul style="list-style-type: none"><li>• Memory Allocation / De-allocation overhead</li><li>• Memory lifecycle tracking overhead</li><li>• Interoperability</li></ul>	<ul style="list-style-type: none"><li>• Message Allocation Pools?</li><li>• Joint TSB/PDU block base class</li></ul>

# Measuring Overhead in Burst Systems

- ▶ Start with several simple flowgraphs (~1.8 femtobalints each)
- ▶ Let them run unthrottled and measure max throughput and latency
- ▶ Both use the “same” work function for pdus and message ports
- ▶ Warning: These were conducted with 100-200 item packets, there are lots of variables in this benchmark and this represents one test ... not a conclusive result yet.



TSB 0 Blocks (B0)




PDU 0 Blocks (B0)



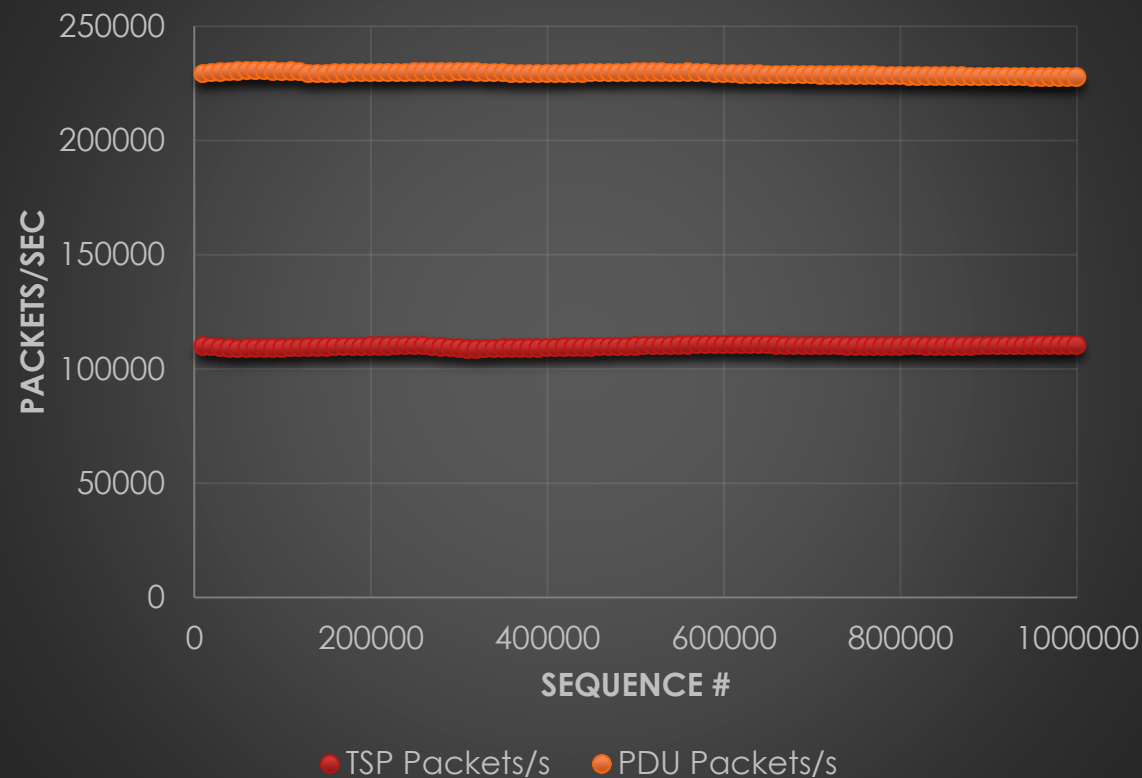
# Measuring Overhead in Burst Systems

## B0 Latency Comparison



TSB 0 Blocks (B0)

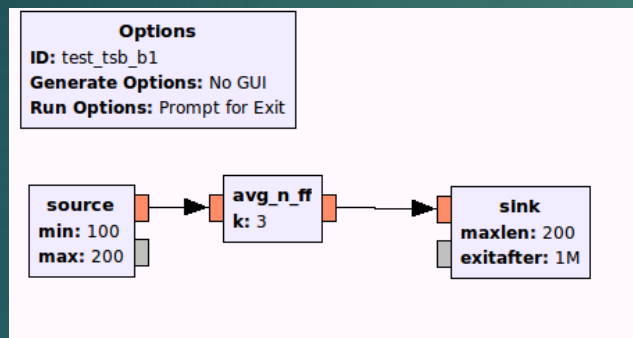
## B0 Packet Throughput



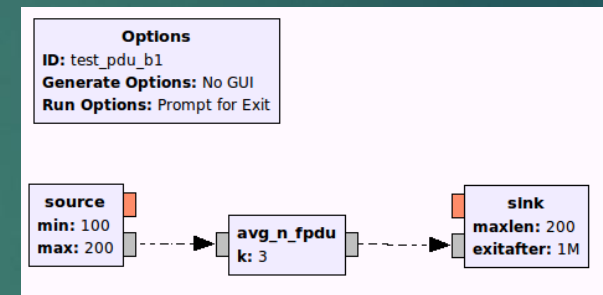
PDU 0 Blocks (B0)

# Measuring Overhead in Burst Systems

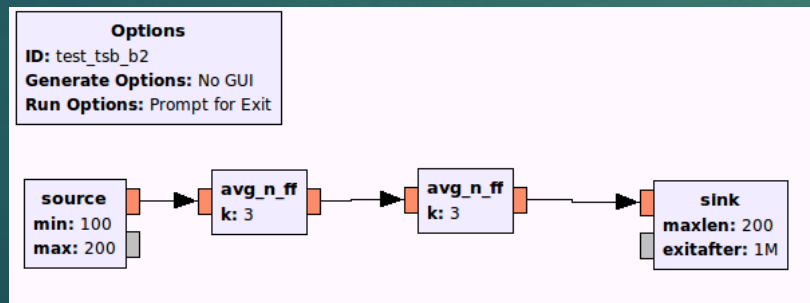
- ▶ Similar B1 and B2 tests, measure latency and throughput through simple burst processing blocks (function is a block average which shortens packets by averaging many little chunks together in this case ... operation is identical for PDU and TSB implementations.)



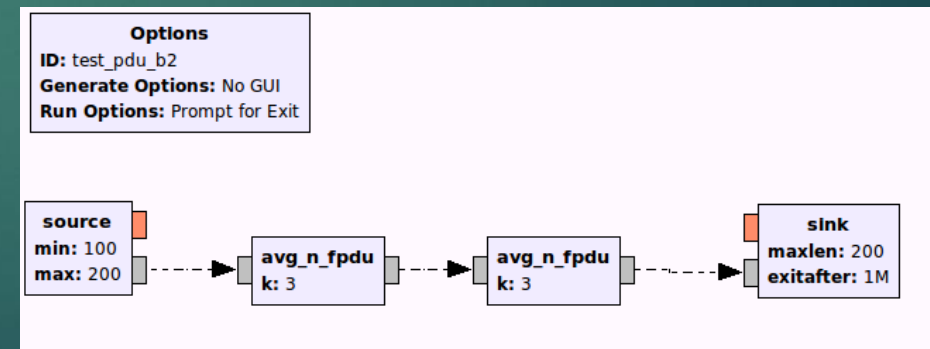
TSB 1 Blocks (B1)



PDU 1 Blocks (B1)




TSB 2 Blocks (B2)




PDU 2 Blocks (B2)

# Measuring Overhead in Burst Systems


## B1 Latency




## B1 Throughput



## B2 Latency



## B2 Throughput



# Measuring Overhead in Burst Systems

## ► Conclusions:

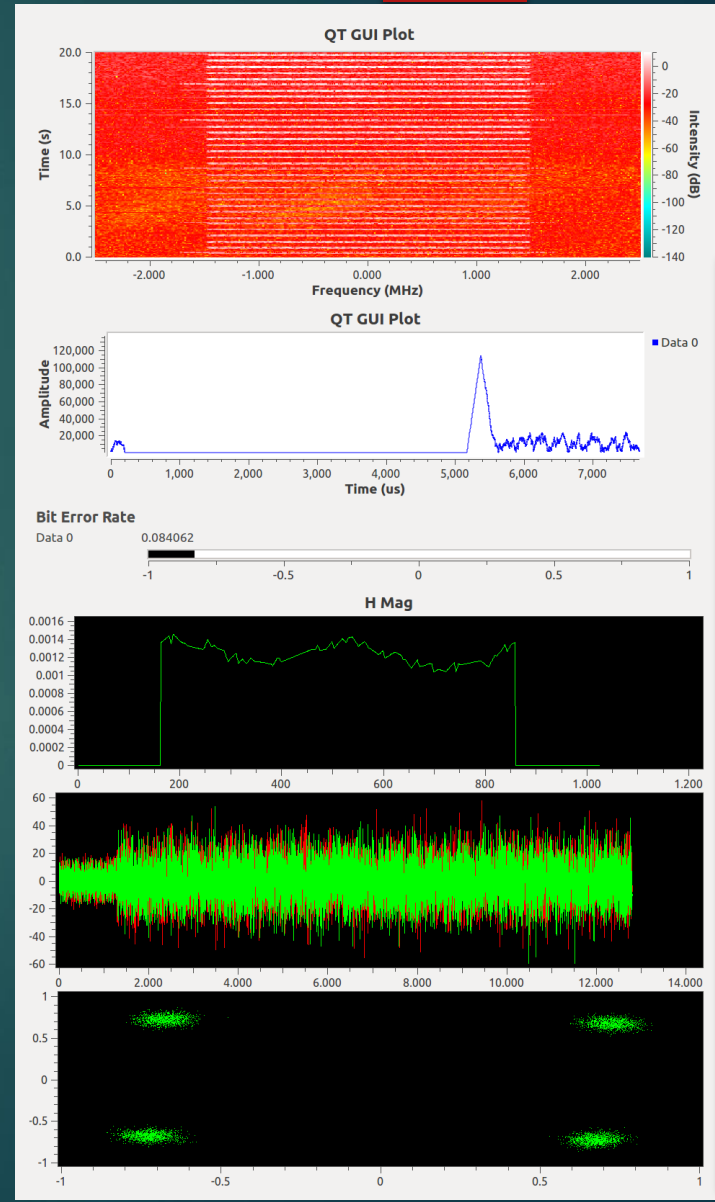
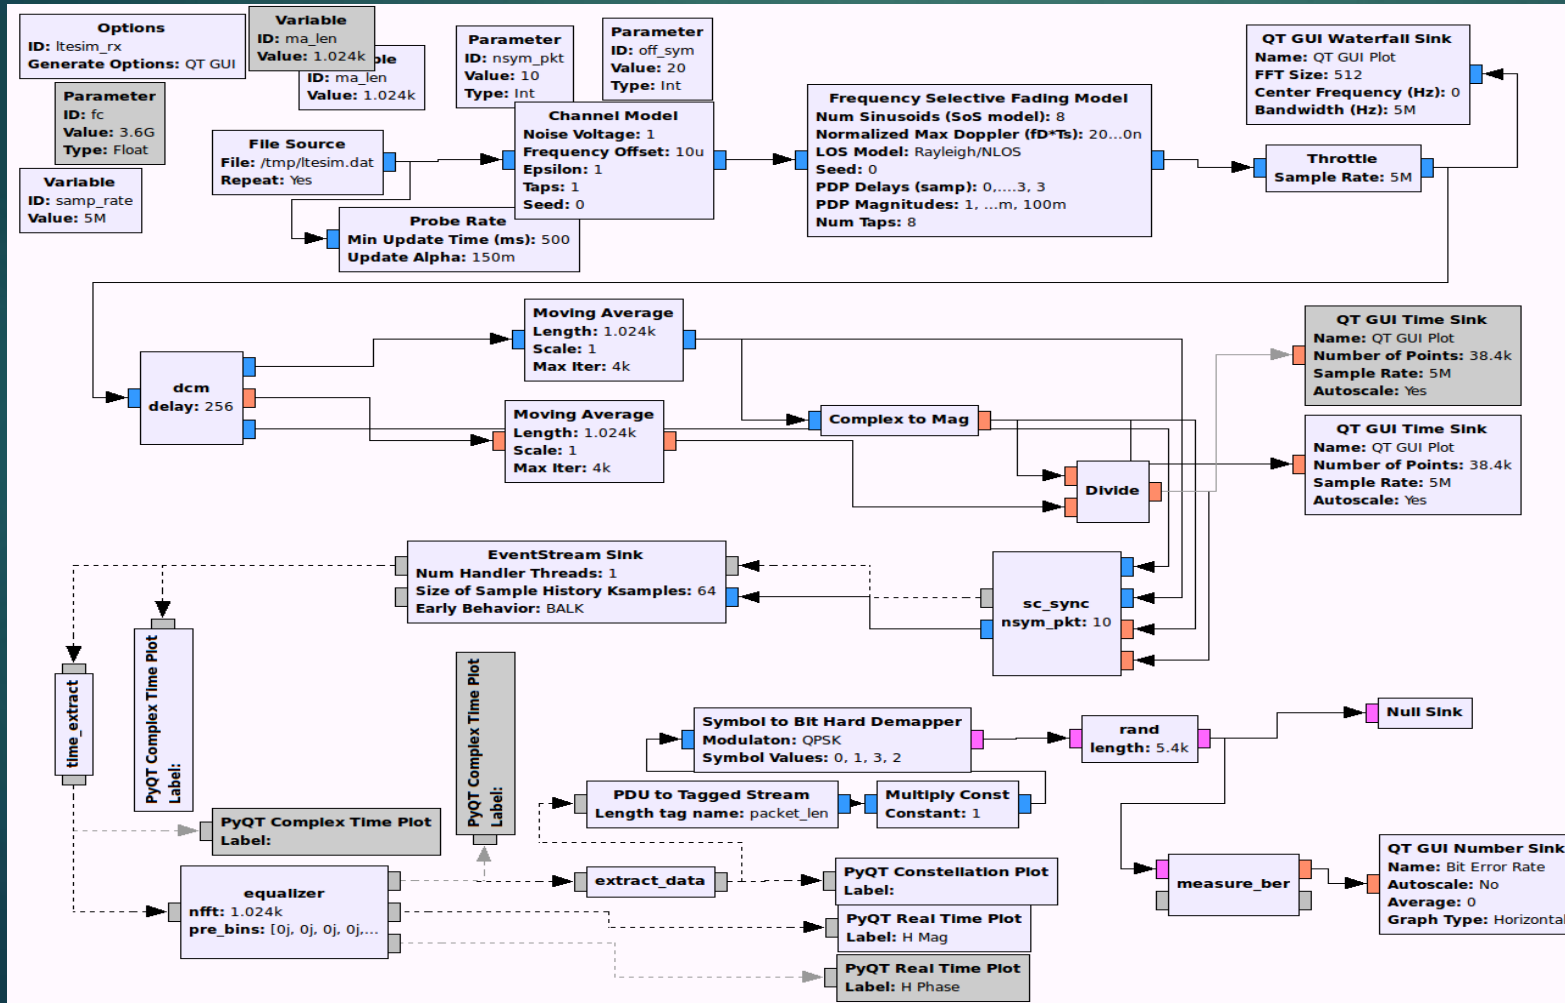
- More measurement is required
- PDU performance seems to be far better than some have speculated ...

## ► Next steps

- Third party verification (code is on github @ <https://github.com/osh/gr-chunky>)
- Automate and standardize measurement process
- Ensure all blocks are best-case optimized
- Common base class for TSB/PDU processing blocks
- Measurement on alternative platforms (i.e. Zynq)

# GR-OFDMA Message Based Modem

- ▶ Now lets build some modems with message connections!
- ▶ Almost posted on github ... soon





# http://pybombs.info



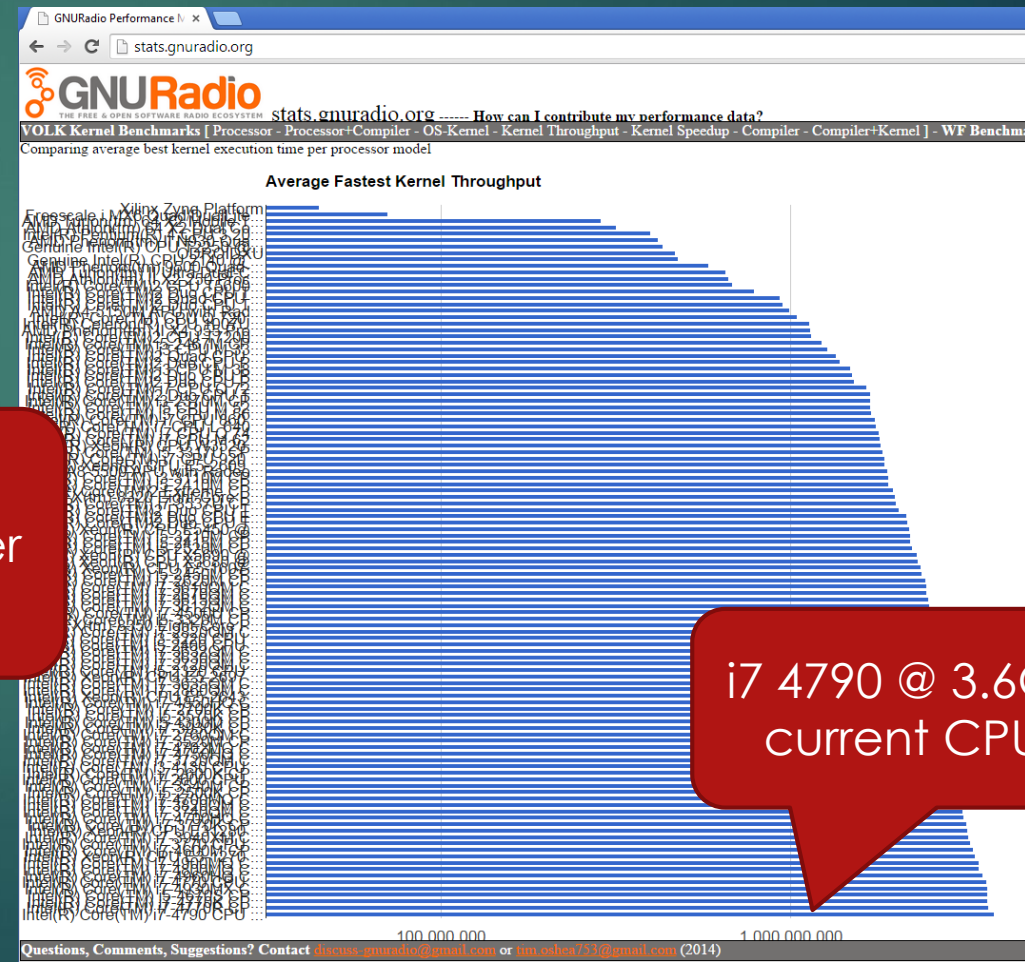

Handy new URL ...

Fancier “app store”  
Does anyone actually use this?

Please send github pull requests for  
new recipes to add!

# <http://stats.gnuradio.org>

- ▶ Summarized VOLK kernel and waveform benchmark performance
- ▶ Processing performance comparisons
- ▶ Needs some help from someone good with web visualization



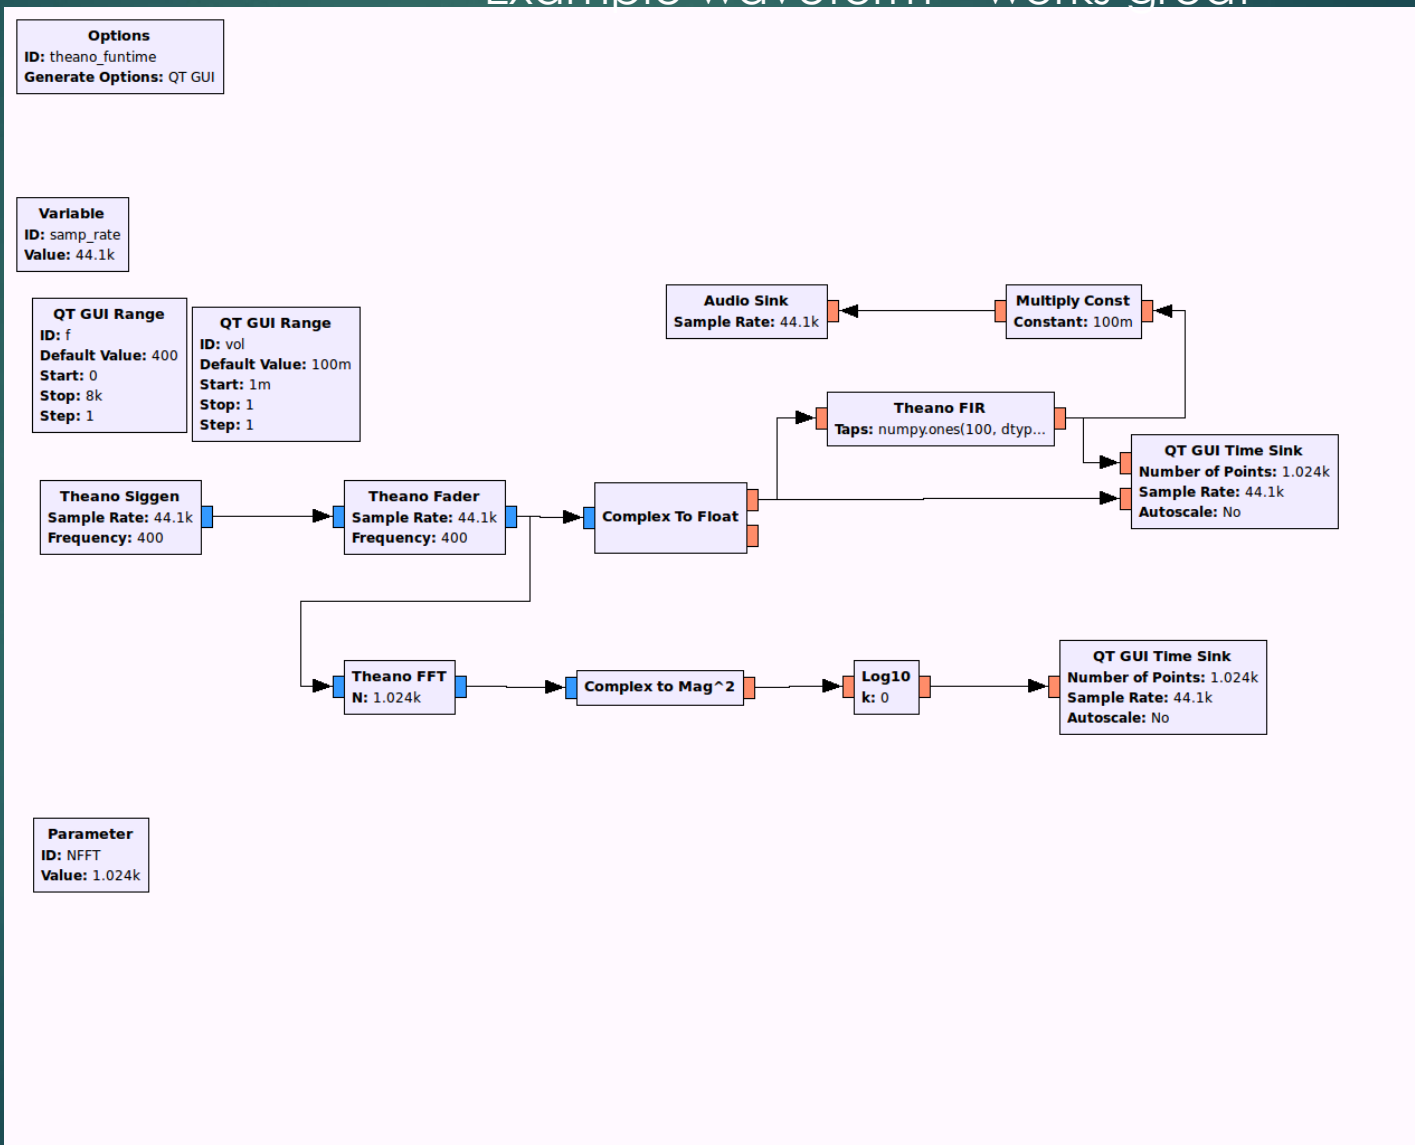


# gr-theano

- ▶ Theano is a stand alone python library used for accelerating large matrix mathematical calculations on GPU cards
- ▶ Heavily used in the Machine Learning community
- ▶ OpenCL and CUDA backends (CUDA is the most mature for now)
- ▶ Optimizes numpy-like python vector operations and compiles them to kernels for you, manages the data transport for you
- ▶ Great for building a monolithic algorithm block – still problematic for chaining gpu-blocks together
- ▶ Makes writing python-based gpu blocks insanely easy
- ▶ Online benchmarks claim 20x speedup on GPU for some algorithms
- ▶ Gr-Theano still needs benchmarking ...
- ▶ Gr-Theano is on github @ <https://github.com/osh/gr-theano>

# gr-theano

Example waveform – works great



# gr-theano

Two handy example blocks .... So concise, wow

```
class fir(gr.sync_block):
    x = T.matrix("x")

    def set_taps(self, taps):
        print "set_taps"
        self.b = theano.shared(numpy.vstack([taps]), name="b")
        self.set_history(taps.size)

    def __init__(self, taps):
        gr.sync_block.__init__(self,
                                name="theano_fir",
                                in_sig=[numpy.float32],
                                out_sig=[numpy.float32])
        self.set_taps(taps)

    self.f = theano.function(
        inputs = [self.x],
        outputs=[T.signal.conv.conv2d(self.x,self.b)],
        updates={},
        name ="f")

    def work(self, input_items, output_items):
        out = output_items[0]
        o = self.f( numpy.vstack([ input_items[0] ] ) );
        out[:] = o[0][0,:];
        return len(output_items[0])
```

```
class fft(gr.sync_block):
    def __init__(self, N):
        gr.sync_block.__init__(self,
                                name="theano_fft",
                                in_sig=[numpy.complex64],
                                out_sig=[numpy.complex64])

        self.set_output_multiple(N);
        self.N = N
        x = T.cmatrix("x")
        w = theano.shared(numpy.ones(self.N,
dtype="complex64"), name="w")
        self.f = theano.function(
            inputs=[x],
            outputs=[T.fourier.fft(x*w, n=N, axis=1)],
            updates={},
            name = "f")

    def work(self, input_items, output_items):
        n = len(input_items[0])/self.N
        for i in range(0,n):
            inmat =
numpy.vstack([input_items[0][self.N*i:self.N*(i+1)]]);
            omat = self.f(inmat);
            output_items[0][self.N*i:self.N*(i+1)] = omat[0];
        return len(output_items[0])
```

That's it for this year

